

Bayesian workflow for disease transmission modeling in Stan

Eustat – XXXIII International Statistical Seminar

Julien Riou, MD PhD

Institute of Social and Preventive Medicine, University of Bern, Switzerland

Preface

- Objective: fit transmission models in Stan
- Based on Grinsztajn et al., 2020 ([link](#))
- Prerequisites:
 - general understanding of Bayesian inference
 - basic programming with R and Stan
- All material is available on
`https://github.com/jriou/bayesian_workflow`

Outline

- **Introduction**
- (Quick notice: Bayesian inference with Stan)
- Fitting a simple SIR
- Simulations to understand the model
- Scaling up ODE-based models
- Extensions

Introduction

Models of disease transmission:

- Interpretability: **mechanistic**, phenomenological
- Scale: agent-based, **population-based**
- Framework: **deterministic**, stochastic
- Data-generating mechanisms: incubation, contagion, immunity...

Introduction

Models of disease transmission:

- Interpretability: **mechanistic**, phenomenological
- Scale: agent-based, **population-based**
- Framework: **deterministic**, stochastic
- Data-generating mechanisms: incubation, contagion, immunity...

Mechanistic + population-based + deterministic

→ **ordinary differential equations (ODE)-based compartmental model**

Introduction

ODE-based compartmental model:


- Divide the population into homogeneous groups (**compartments**)
- Define the **flows** between compartments with ODEs
- Define initial conditions (at t_0)
- Solve for the time-dependent volume in each compartment

Introduction

ODE-based compartmental model:

- Divide the population into homogeneous groups (**compartments**)
- Define the **flows** between compartments with ODEs
- Define initial conditions (at t_0)
- Solve for the time-dependent volume in each compartment

The **susceptible-infectious-recovered** (SIR) model:

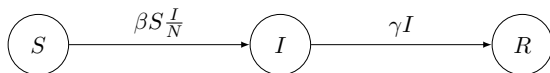


$$\frac{dS}{dt} = -\beta S \frac{I}{N}$$

$$\frac{dI}{dt} = \beta S \frac{I}{N} - \gamma I$$

$$\frac{dR}{dt} = \gamma I$$

Introduction



$$\frac{dS}{dt} = -\beta S \frac{I}{N}$$

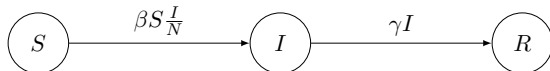
$$\frac{dI}{dt} = \beta S \frac{I}{N} - \gamma I$$

$$\frac{dR}{dt} = \gamma I$$

Where:

- $S(t)$ is the number of people **susceptible** to infection
- $I(t)$ is the number of people **infected** (i.e. the prevalence)
- $R(t)$ is the number of people **recovered** (lifelong immunity)
- N is the population size ($S(t) + I(t) + R(t) = N$ for any t)
- β is the **infectious contact rate** (per day per person)
- γ is the **recovery rate** (1/infectious period)

Introduction



$$\frac{dS}{dt} = -\beta S \frac{I}{N}$$

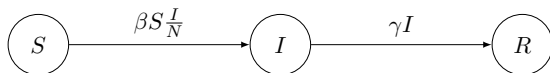
$$\frac{dI}{dt} = \beta S \frac{I}{N} - \gamma I$$

$$\frac{dR}{dt} = \gamma I$$

Intuition behind the SIR model:

- $I(t)/N$ is the proportion of infected (and infectious)
- $\beta I(t)/N$ is the daily number of contacts with infectious people
- hence each day, $\beta S I(t)/N$ people become infected (the **force of infection**)

Introduction



$$\frac{dS}{dt} = -\beta S \frac{I}{N}$$

$$\frac{dI}{dt} = \beta S \frac{I}{N} - \gamma I$$

$$\frac{dR}{dt} = \gamma I$$

Assumptions behind the SIR model:

- homogeneous mixing
- β and γ constant over time
- all infections are observed
- no incubation, exponentially-distributed recovery
- lifelong immunity
- stable population

Introduction

Simulate in R with package `deSolve`:

- set compartments and differential equations

```
> ## Set model ----  
> seir = function(t, x, parms, ...) {  
+   with(as.list(c(parms, x)), {  
+     dS = - beta*S*I/(S+I+R)  
+     dI = beta*S*I/(S+I+R) - gamma*I  
+     dR = gamma*I  
+     list(c(dS, dI, dR))  
+   })  
+ }
```

Introduction

Simulate in R with package deSolve:

- set compartments and differential equations

```
> ## Set model ----  
> seir = function(t, x, parms, ...) {  
+   with(as.list(c(parms, x)), {  
+     dS = - beta*S*I/(S+I+R)  
+     dI = beta*S*I/(S+I+R) - gamma*I  
+     dR = gamma*I  
+     list(c(dS, dI, dR))  
+   })  
+ }
```

- set (fixed) values for $\beta = 0.8$; $\rho = 1/7$; $S_0 = 100,000 - 50$; $I_0 = 50$ and $R_0 = 0$

```
> ## Set parameters ----  
> pars = c(beta = 0.8,  
+          gamma = 1/7  
+ )
```

```
> ## Set initial values ----  
> N_0 = 100000  
> I_0 = 50  
> inits = c(  
+   S = N_0 - I_0,  
+   I = I_0,  
+   R = 0  
+ )
```

Introduction


- solve the ODE system numerically (Runge-Kutta 4th order) to obtain unique solutions for $S(t)$, $I(t)$ and $R(t)$

$$f(\beta, \gamma, S_0, I_0, R_0) = \{S(t), I(t), R(t)\}$$

```
> ## Simulate ----
> times = seq(0,50,by=1)
> sim_data = ode(inits, times, seir, pars,method="rk4")
> tibble(sim_data)
# A tibble: 51 x 1
  sim_data[,"time"] [,"S"] [,"I"] [,"R"]
      <dbl>      <dbl> <dbl> <dbl>
1           0  99950      50      0
2           1  99894.    96.3    10.1
3           2  99785.   186.    29.5
4           3  99576.   357.    66.9
5           4  99176.   685.   139.
6           5  98415.  1308.   276.
7           6  96984.  2478.   538.
8           7  94350.  4621.  1030.
9           8  89692.  8374.  1934.
10          9  82009. 14457. 3533.
# ... with 41 more rows
```


Introduction

with $\beta = 0.8$; $\rho = 1/7$; $S_0 = 100000 - 50$; $I_0 = 50$ and $R_0 = 0$



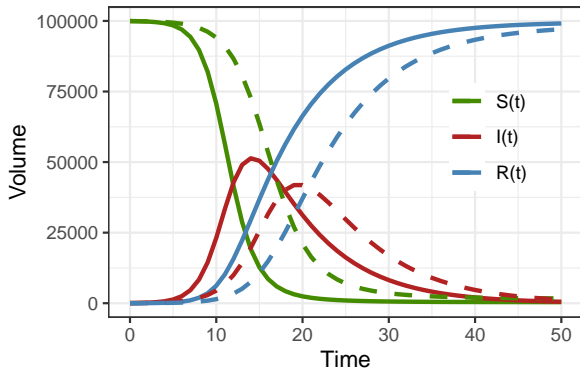
Introduction

with $\beta = 1.1$ instead of 0.8, we get



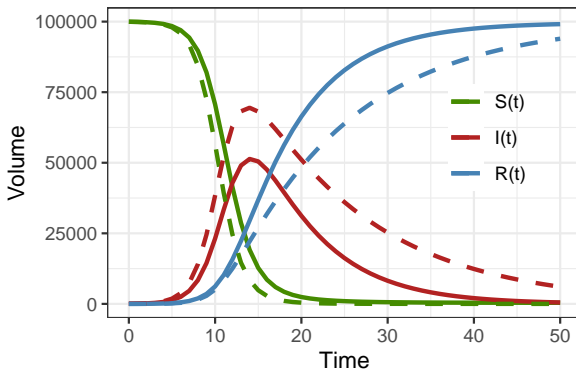
Introduction

with $\beta = 0.6$ instead of 0.8, we get




Introduction

with $\gamma = 1/14$ instead of $1/7$, we get



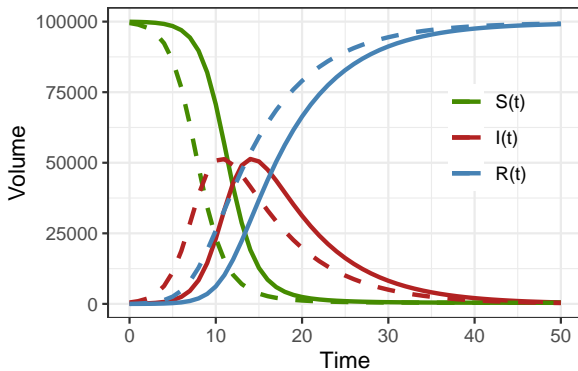
Introduction

with $\gamma = 1/4$ instead of $1/7$, we get



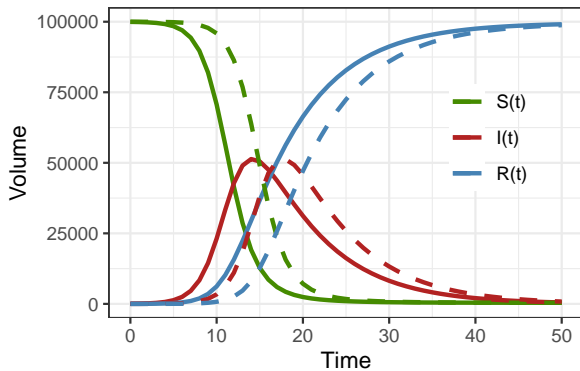
Introduction

with $I(0) = 500$ instead of 50, we get



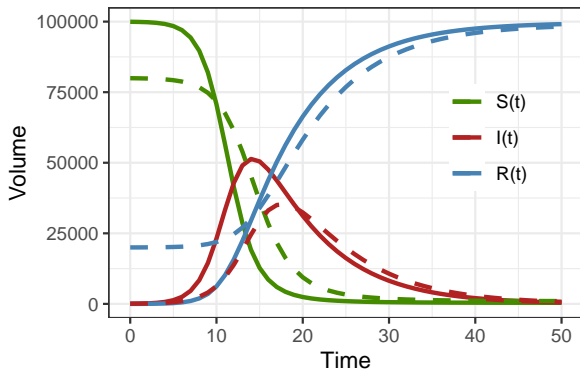
Introduction

with $I(0) = 5$ instead of 50, we get



Introduction

with $R(0) = 20,000$ instead of 0, we get



Introduction

Compartmental models have many uses:

- formalize and put numerical values on **general concepts** (herd immunity, vaccination threshold...)
- get **mechanistic insight** about an epidemic (transmissibility levels, drivers of transmission)

$$\mathcal{R}_0 = \frac{\beta}{\gamma}$$

- produce precise **forecasts** (based on mechanisms)

Introduction

Compartmental models have many uses:

- formalize and put numerical values on **general concepts** (herd immunity, vaccination threshold...)
- get **mechanistic insight** about an epidemic (transmissibility levels, drivers of transmission)

$$\mathcal{R}_0 = \frac{\beta}{\gamma}$$

- produce precise **forecasts** (based on mechanisms)

→ all these uses are based on **numerical values** for β , ρ and the initial conditions and their **uncertainty**

Introduction

Enters **Bayesian inference**:

- infer parameter values by **integrating data and prior knowledge**
- more efficient for complex models (high dimensionality)
- rigorously quantify and propagate uncertainty in parameter estimates and forecast

Introduction

Enters **Bayesian inference**:

- infer parameter values by **integrating data and prior knowledge**
- more efficient for complex models (high dimensionality)
- rigorously quantify and propagate uncertainty in parameter estimates and forecast

→ Markov Chain Monte Carlo (MCMC) methods and **Stan**

Outline

- Introduction
- **(Quick notice: Bayesian inference with Stan)**
- Fitting a simple SIR
- Simulations to understand the model
- Scaling up ODE-based models
- Extensions

(Bayesian inference with Stan)

General principle of Bayesian inference:

- specify a complete Bayesian model
 - consider data $y = \{y_1, \dots, y_n\}$ and parameter θ
 - specify an **observation model**

$$\Pr(y|\theta) = \prod_n \text{normal}(y_n|\theta, 1)$$

- complete the model with a **prior distribution**

$$\Pr(\theta) = \text{normal}(0, 1)$$

- sample the **posterior distribution** of the parameter

(Bayesian inference with Stan)

Stan is a probabilistic programming framework for Bayesian inference

- it is designed to let the user **focus on modeling** while inference happens under the hood
- object-oriented language (based on C++) that supports many operations, probability densities and ODE solvers
- extremely **efficient** MCMC algorithm (Hamiltonian Monte Carlo)
- **diagnostic tools** to evaluate the inference
- interfaces in R (package `rstan`), python, julia...

(Bayesian inference with Stan)

Programming in Stan is structured in **blocks**:

- the data block defines data variables

```
data {  
  int N;  
  real y[N];  
}
```

- the parameters block defines parameters

```
parameters {  
  real theta;  
}
```

- the model block defines the **target log probability density function**

```
model {  
  theta ~ normal(0,1);  
  y ~ normal(theta,1);  
}
```

- save in model_linear.stan

(Bayesian inference with Stan)

We then explore the target with Stan's MCMC **sampler**:

- load `rstan` package

```
## Setup ----  
library(rstan)  
options(mc.cores = parallel::detectCores())
```

- simulate $N = 50$ data points with $\theta = 0.7$


```
## Simulate data ----  
N = 50  
theta = 0.7  
y = rnorm(N,theta,1)  
input_data = list(N=N,y=y)
```

- run MCMC sampling

```
## Sample ----  
fit = stan(file='model_linear.stan',  
           data=input_data,  
           chains=4,  
           iter=1000)
```


(Bayesian inference with Stan)

We use **multiple chains** that should converge after warm-up




(Bayesian inference with Stan)

We use **multiple chains** that should converge after warm-up



(Bayesian inference with Stan)

The post-warm-up samples of θ approximate its **posterior distribution**



(Bayesian inference with Stan)

We run **basic diagnosis tools**: divergences, tree depth, energy

```
> check_hmc_diagnostics(fit)
```

```
Divergences:
```

```
0 of 2000 iterations ended with a divergence.
```

```
Tree depth:
```

```
0 of 2000 iterations saturated the maximum tree depth of 10.
```

```
Energy:
```

```
E-BFMI indicated no pathological behavior.
```

(Bayesian inference with Stan)

Printing the object gives:

- **diagnostics**: effective sample size, Gelman-Rubin \hat{R}
- **inference**: full posterior distribution of θ

```
> print(fit)
Inference for Stan model: model_linear.
4 chains, each with iter=1000; warmup=500; thin=1;
post-warmup draws per chain=500, total post-warmup draws=2000.

      mean se_mean  sd  2.5%   25%   50%   75%  97.5% n_eff Rhat
theta  0.78    0.01 0.14   0.51   0.69   0.78   0.87   1.05   760    1
lp__   -18.30    0.03 0.73 -20.28 -18.44 -18.03 -17.86 -17.81   799    1

Samples were drawn using NUTS(diag_e) at Thu Nov 12 19:15:33 2020.
For each parameter, n_eff is a crude measure of effective sample size,
and Rhat is the potential scale reduction factor on split chains (at
convergence, Rhat=1).
```


Outline

- Introduction
- (Quick notice: Bayesian inference with Stan)
- **Fitting a simple SIR**
- Simulations to understand the model
- Scaling up ODE-based models
- Extensions

Fitting a simple SIR

Example data: outbreak of influenza A (H1N1) at a **British boarding school** in 1978 (available in R package `outbreaks`)

- 763 students, 512 had symptoms
- daily number of students in bed over 14 days (**prevalence** data)



Fitting a simple SIR

Specifying the model:

- prevalence data: \mathbb{I}_t with $t \in \{1, \dots, 14\}$
- parameters to estimate: $\theta = \{\beta, \gamma, \phi\}$
- parameters that will remain fixed: $\{S_0 = 762, I_0 = 1, R_0 = 0\}$
- map data \mathbb{I}_t to SIR model output $I(t)$ using an observation model with an appropriate **probability distribution**:

$$\Pr(\mathbb{I}|\theta) = \prod_{t=1}^{14} \text{neg-bin}(\mathbb{I}_t | I(t), \phi)$$

- **prior distributions**

$$\Pr(\beta) = \text{exponential}(1)$$

$$\Pr(1/\gamma) = \text{normal}(2, 0.5)$$

$$\Pr(1/\phi) = \text{exponential}(5)$$

Fitting a simple SIR

We define the ODE system in the function block

```
functions {  
  real[] sir(real t, real[] y, real[] theta, real[] x_r, int[] x_i) {  
  
    real S = y[1];  
    real I = y[2];  
    real R = y[3];  
    real N = x_i[1];  
  
    real beta = theta[1];  
    real gamma = theta[2];  
  
    real dS_dt = -beta * I * S / N;  
    real dI_dt = beta * I * S / N - gamma * I;  
    real dR_dt = gamma * I;  
  
    return {dS_dt, dI_dt, dR_dt};  
  }  
}
```

⚠ Be careful of the signature and formats!

Fitting a simple SIR

We declare the data variables in the data block

```
data {  
  int<lower=1> T;  
  real y0[3];  
  real t0;  
  real ts[T];  
  int N;  
  int cases[T];  
}
```


Fitting a simple SIR

We declare the data variables in the data block

```
data {  
  int<lower=1> T;  
  real y0[3];  
  real t0;  
  real ts[T];  
  int N;  
  int cases[T];  
}
```

and define additional data variables in transformed data

```
transformed data {  
  real x_r[0];  
  int x_i[1];  
  x_i[1]=N;  
}
```

Fitting a simple SIR

Similarly, parameters are declared in the `parameters` block

```
parameters {  
  real<lower=0> beta;  
  real<lower=0> recovery_time;  
  real<lower=0> phi_inv;  
}
```

⚠ It sometimes makes more sense to transform some parameters (e.g., recovery rate γ and overdispersion ϕ) to improve interpretability

Fitting a simple SIR

In transformed parameters, we define additional parameters and solve the ODE system

```
transformed parameters{  
  real y[T,3];  
  real phi = 1. / phi_inv;  
  real gamma = 1. / recovery_time;  
  real theta[2];  
  theta[1] = beta;  
  theta[2] = gamma;  
  
  y = integrate_ode_rk45(sir, y0, t0, ts, theta, x_r, x_i);  
}
```

Fitting a simple SIR

```
y = integrate_ode_rk45(sir, y0, t0, ts, theta, x_r, x_i);
```

Two crucial points:

- be careful about the **formats and signatures**
 - the ODE output `y` is an array of size $T \times 3$ (number of time steps and number of compartments)
 - `sir` is the name of the function defined in the `function` block
 - `y0` is an array of size 3 defined in the `data` block
 - `ts` is an array of size T defined in the `data` block
 - `theta` is an array of size 2 storing the parameters
 - `x_r` is defined as empty in transformed data, but can be used to store fixed real values
 - `x_i` is an array of size 1 storing the population size N (can also be used to store fixed integer values)

Fitting a simple SIR

```
y = integrate_ode_rk45(sir, y0, t0, ts, theta, x_r, x_i);
```

Two crucial points:

- two ODE solvers are available:
 - `integrate_ode_rk45` uses the Runge-Kutta method (quicker but non-adapted to stiff systems)
 - `integrate_ode_bdf` uses the backward differentiation method (slower but adapted to stiff systems)

Fitting a simple SIR

In the `model` block, we write the priors and the observation model

```
model {  
  // priors  
  beta ~ exponential(1);  
  recovery_time ~ normal(2,0.5);  
  phi_inv ~ exponential(5);  
  
  // observation model  
  cases ~ neg_binomial_2(col(to_matrix(y),2), phi);  
}
```

⚠ It's important that the chosen distributions correspond with the boundaries set in the parameters block (`<lower=0>`)

⚠ `col(to_to_matrix(y))` extracts the 2nd column of `y`

Fitting a simple SIR

Last, we add a `generated quantities` block that does not influence sampling and can be used for “post-processing”:

- reproduction number $\mathcal{R}_0 = \beta/\gamma$
- model predictions of prevalence from the negative binomial

```
generated quantities {  
  real R0 = beta/gamma;  
  real pred_cases[T];  
  pred_cases = neg_binomial_2_rng(col(to_matrix(y),2), phi);  
}
```

Fitting a simple SIR

In summary:

- `functions`: define the ODE system (⚠ signature and formats)
- `data`: declare data variables that will be provided
- `transformed data`: additional quantities that can be computed internally or from data variables
- `parameters`: declare parameters (⚠ boundaries)
- `transformed parameters`: quantities that can be computed internally or from data or parameters variables, including the ODE output (⚠ signature and format)
- `model`: priors and observation model
- `generated quantities`: additional quantities that can be computed without influencing the sampling

Fitting a simple SIR

As before, we conduct the inference from R with the package `rstan`:

```
## Format input ----  
# prevalence data  
cases = influenza_england_1978_school$in_bed  
N = 763  
n_days = 14  
t0 = 0  
t = 1:n_days  
  
# initial conditions  
i0 = 1  
s0 = N - i0  
r0 = 0  
y0 = c(s0, i0, r0)  
  
# put into list  
input_data = list(T = n_days, y0 = y0, t0 = t0, ts = t, N = N, cases = cases)
```

⚠ data is put in a list with names matching the data block in Stan

Fitting a simple SIR

Hit the inference button!

```
## Sample ----  
fit = stan(file='sir_negbin.stan',  
           data=input_data,  
           chains=4,  
           iter=1000)
```

Fitting a simple SIR

Run basic diagnosis tools:

```
> check_hmc_diagnostics(fit)
```

Divergences:

0 of 2000 iterations ended with a divergence.

Tree depth:


0 of 2000 iterations saturated the maximum tree depth of 10.

Energy:

E-BFMI indicated no pathological behavior.


Fitting a simple SIR

Examine trace plots:




Fitting a simple SIR

Examine chain mixing:



Fitting a simple SIR

Posterior predictive check:



Fitting a simple SIR

Print the results:

```
> print(fit,pars=c("beta","gamma","phi","R0","recovery_time"))
Inference for Stan model: sir_negbin.
4 chains, each with iter=1000; warmup=500; thin=1;
post-warmup draws per chain=500, total post-warmup draws=2000.
```

	mean	se_mean	sd	2.5%	25%	50%	75%	97.5%	n_eff	Rhat
beta	1.73	0.00	0.06	1.61	1.70	1.73	1.76	1.85	1049	1
gamma	0.53	0.00	0.04	0.44	0.50	0.53	0.56	0.63	1382	1
phi	9.61	0.22	6.13	2.94	5.72	8.31	11.80	23.40	743	1
R0	3.27	0.01	0.29	2.79	3.09	3.25	3.42	3.96	1403	1
recovery_time	1.89	0.00	0.16	1.59	1.79	1.88	1.98	2.25	1410	1

Samples were drawn using NUTS(diag_e) at Thu Nov 12 19:10:02 2020.
For each parameter, n_eff is a crude measure of effective sample size,
and Rhat is the potential scale reduction factor on split chains (at
convergence, Rhat=1).

Outline

- Introduction
- (Quick notice: Bayesian inference with Stan)
- Fitting a simple SIR
- **Simulations to understand the model**
- Scaling up ODE-based models
- Extensions

Simulations to understand the model

Acknowledgements & ressources

- Michael Betancourt's *Introduction to Stan*
https://betanalpha.github.io/assets/case_studies/stan_intro.html
- Daniel Lee's *ODEs in Stan*
https://youtu.be/hJ34_xJhYeY
- Richard McElreath's *Statistical rethinking*
<https://youtu.be/4WVe1CswXo4>