

Bayesian workflow for disease transmission modeling in Stan

Eustat – XXXIII International Statistical Seminar

Julien Riou, MD PhD

Institute of Social and Preventive Medicine, University of Bern, Switzerland

Preface

- Objective: fit transmission models in Stan
- Based on Grinsztajn et al., 2020 ([link](#))
- Prerequisites:
 - general understanding of Bayesian inference
 - basic programming with R and Stan
- All material is available on
`https://github.com/jriou/bayesian_workflow`

Outline

- Introduction
- (Quick notice: Bayesian inference with Stan)
- Simple SIR
- Using simulated data
- Scaling up ODE-based models
- Extensions

Introduction

Models of disease transmission:

- Interpretability: mechanistic, phenomenological

Introduction

Models of disease transmission:

- Interpretability: mechanistic, phenomenological
- Data-generating mechanisms: incubation, contagion, immunity...

Introduction

Models of disease transmission:

- Interpretability: mechanistic, phenomenological
- Data-generating mechanisms: incubation, contagion, immunity...
- Scale: agent-based, population-based

Introduction

Models of disease transmission:

- Interpretability: mechanistic, phenomenological
- Data-generating mechanisms: incubation, contagion, immunity...
- Scale: agent-based, population-based
- Framework: deterministic, stochastic

Introduction

Models of disease transmission:

- Interpretability: mechanistic, phenomenological
- Data-generating mechanisms: incubation, contagion, immunity...
- Scale: agent-based, population-based
- Framework: deterministic, stochastic

Mechanistic + population-based + deterministic

→ ODE-based compartmental model (e.g., SIR)

Introduction

ODE-based compartmental model:

- Divide the population into homogeneous groups (compartments)

Introduction

ODE-based compartmental model:

- Divide the population into homogeneous groups (compartments)
- Define flows between compartments with differential equations

Introduction

ODE-based compartmental model:

- Divide the population into homogeneous groups (compartments)
- Define flows between compartments with differential equations
- Define initial conditions

Introduction


ODE-based compartmental model:

- Divide the population into homogeneous groups (compartments)
- Define flows between compartments with differential equations
- Define initial conditions
- Solve for the time-dependent volume in each compartment

Introduction

ODE-based compartmental model:

- Divide the population into homogeneous groups (compartments)
- Define flows between compartments with differential equations
- Define initial conditions
- Solve for the time-dependent volume in each compartment



$$\frac{dS}{dt} = -\beta S \frac{I}{N}$$

$$\frac{dI}{dt} = \beta S \frac{I}{N} - \gamma I$$

$$\frac{dR}{dt} = \gamma I$$

Introduction

Simulate in R with package `deSolve`:

- set compartments and differential equations

```
> ## Set model ----  
> seir = function(t, x, parms, ...) {  
+   with(as.list(c(parms, x)), {  
+     dS = - beta*S*I/(S+I+R)  
+     dI = beta*S*I/(S+I+R) - gamma*I  
+     dR = gamma*I  
+     list(c(dS, dI, dR))  
+   })  
+ }
```

Introduction

Simulate in R with package deSolve:

- set compartments and differential equations

```
> ## Set model ----
> seir = function(t, x, parms, ...) {
+   with(as.list(c(parms, x)), {
+     dS = - beta*S*I/(S+I+R)
+     dI = beta*S*I/(S+I+R) - gamma*I
+     dR = gamma*I
+     list(c(dS, dI, dR))
+   })
+ }
```

- set (fixed) values for β , ρ and initial conditions

$$\beta = 0.8; \gamma = 1/7; S(0) = 100,000 - 50; I(0) = 50; R(0) = 0$$

```
> ## Set parameters ----
> pars = c(beta = 0.8,
+   gamma = 1/7
+ )
```

```
> ## Set initial values ----
> N_0 = 100000
> I_0 = 50
> inits = c(
+   S = N_0 - I_0,
+   I = I_0,
+   R = 0
+ )
```


Introduction

- solve the ODE system numerically (Runge-Kutta 4th order)

```
> ## Simulate ----
> times = seq(0,50,by=1)
> sim_data = ode(inits, times, seir, pars,method="rk4")
> tibble(sim_data)
# A tibble: 51 x 1
  sim_data["time"] [, "S"] [, "I"] [, "R"]
      <dbl>      <dbl>      <dbl>      <dbl>
1           0 99950         50         0
2           1 99894.        96.3       10.1
3           2 99785.       186.        29.5
4           3 99576.       357.        66.9
5           4 99176.       685.       139.
6           5 98415.      1308.       276.
7           6 96984.      2478.       538.
8           7 94350.      4621.      1030.
9           8 89692.      8374.      1934.
10          9 82009.     14457.     3533.
# ... with 41 more rows
```


Introduction

- we obtain (deterministic) values for $S(t)$, $I(t)$ and $R(t)$




- these quantities have real-world interpretations (respectively susceptibility, prevalence, and cumulative attack rate)

Introduction

- the result is entirely determined by the chosen values for β , ρ and the initial conditions


with $\beta = 1.1$ instead of 0.8, we get



Introduction

- the result is entirely determined by the chosen values for β , ρ and the initial conditions


with $\beta = 0.6$ instead of 0.8, we get



Introduction

- the result is entirely determined by the chosen values for β , ρ and the initial conditions


with $\gamma = 1/14$ instead of $1/7$, we get



Introduction

- the result is entirely determined by the chosen values for β , ρ and the initial conditions


with $\gamma = 1/4$ instead of $1/7$, we get



Introduction

- the result is entirely determined by the chosen values for β , ρ and the initial conditions


with $I(0) = 500$ instead of 50, we get



Introduction

- the result is entirely determined by the chosen values for β , ρ and the initial conditions


with $I(0) = 5$ instead of 50, we get



Introduction

- the result is entirely determined by the chosen values for β , ρ and the initial conditions

with $R(0) = 20,000$ instead of 0, we get



Introduction

Compartmental models have many uses:

- formalize and quantify general concepts (herd immunity, vaccination threshold...)
- get mechanistic insight about an epidemic (\mathcal{R}_0 , \mathcal{R}_t , impact of interventions...)
- produce forecasts (based on mechanisms)

Introduction

Compartmental models have many uses:

- formalize and quantify general concepts (herd immunity, vaccination threshold...)
- get mechanistic insight about an epidemic (\mathcal{R}_0 , \mathcal{R}_t , impact of interventions...)
- produce forecasts (based on mechanisms)

→ based on numerical values for β , ρ and the initial conditions

Introduction

Enters **Bayesian inference**:

- make the best use of information from data
- easily incorporate prior knowledge
- infer the value (posterior) of the parameters
- propagate uncertainty from data and priors

Introduction

Enters **Bayesian inference**:

- make the best use of information from data
- easily incorporate prior knowledge
- infer the value (posterior) of the parameters
- propagate uncertainty from data and priors

→ Markov Chain Monte Carlo methods and **Stan**

(Bayesian inference with Stan)

General principle of Bayesian inference:

- specify a **complete Bayesian model**
 - consider data $y = \{y_1, \dots, y_n\}$ and parameter θ
 - specify an observation model

$$\Pr(y|\theta) = \prod_n \text{normal}(y_n|\theta, 1)$$

- complete the model with a prior on the parameter

$$\Pr(\theta) = \text{normal}(0, 1)$$

- estimate the **joint probability density function** of the model

(Bayesian inference with Stan)

The **joint probability density function** of the model is given by

$$\Pr(y, \theta) = \prod_{n=1}^N \text{normal_pdf}(y_n \mid \theta, 1) \cdot \text{normal_pdf}(\theta \mid 0, 1)$$

or on the log scale

$$\log \Pr(y, \theta) = \sum_{n=1}^N \text{normal_lpdf}(y_n \mid \theta, 1) + \text{normal_lpdf}(\theta \mid 0, 1)$$

(Bayesian inference with Stan)

Programming in Stan is structured in **blocks**:

- the data block defines data variables

```
data {  
  int N;  
  real y[N];  
}
```

- the parameters block defines parameters

```
parameters {  
  real theta;  
}
```

- the model block defines the **target log probability density function**

```
model {  
  target += normal_lpdf(theta | 0, 1);  
  for (n in 1:N)  
    target += normal_lpdf(y[n] | theta, 1);  
}
```

(Bayesian inference with Stan)

We then explore the target with **Hamiltonian Monte Carlo**:

- load rstan package

```
## Setup ----  
library(rstan)  
options(mc.cores = parallel::detectCores())
```

- simulate $N = 50$ data points with $\theta = 0.7$

```
## Simulate data ----  
N = 50  
theta = 0.7  
y = rnorm(N,theta,1)  
input_data = list(N=N,y=y)
```

- run sampling

```
## Sample ----  
fit = stan(file='example_linear/model_linear.stan',  
           data=input_data,  
           chains=4,  
           iter=1000)
```


(Bayesian inference with Stan)

- print results

```
> print(fit)
Inference for Stan model: model_linear.
4 chains, each with iter=1000; warmup=500; thin=1;
post-warmup draws per chain=500, total post-warmup draws=2000.

      mean se_mean  sd  2.5%   25%   50%   75%  97.5% n_eff Rhat
theta  0.67   0.01 0.14  0.39   0.57  0.66  0.76   0.95  510 1.01
lp__ -75.37   0.03 0.76 -77.52 -75.52 -75.09 -74.92 -74.86  544 1.00
```

Samples were drawn using NUTS(diag_e) at Tue Nov 10 17:14:58 2020.
For each parameter, `n_eff` is a crude measure of effective sample size,
and `Rhat` is the potential scale reduction factor on split chains (at
convergence, `Rhat=1`).

- diagnostics: \hat{R} , divergences, tree depth, energy

```
> check_hmc_diagnostics(fit)

Divergences:
0 of 2000 iterations ended with a divergence.

Tree depth:
0 of 2000 iterations saturated the maximum tree depth of 10.

Energy:
E-BFMI indicated no pathological behavior.
```

Acknowledgements & ressources

- Michael Betancourt's *Introduction to Stan*
https://betanalpha.github.io/assets/case_studies/stan_intro.html
- Daniel Lee's *ODEs in Stan*
https://youtu.be/hJ34_xJhYeY
- Richard McElreath's *Statistical rethinking*
<https://youtu.be/4WVe1CswXo4>